

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY  
UNIVERSITY OF TECHNOLOGY  
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



## MACHINE LEARNING

---

Assignment report

# Gender Classification

---

Lecturer: - Nguyễn Đức Dũng  
Student: - Đinh Hoàng Anh 1952553

HO CHI MINH CITY, NOVEMBER 2021



# Contents

<b>1</b>	<b>Problem description</b>	<b>2</b>
1.1	Overview of gender classification . . . . .	2
1.2	Detailed description . . . . .	2
<b>2</b>	<b>Deep Neural Networks</b>	<b>3</b>
2.1	Deep L-Layer Neural Network . . . . .	3
2.2	Forward Propagation . . . . .	3
2.3	Backward Propagation . . . . .	4
2.4	Update Paramters . . . . .	5
2.5	Cost Function . . . . .	5
<b>3</b>	<b>Dataset</b>	<b>5</b>
<b>4</b>	<b>Implementation</b>	<b>6</b>
4.1	Packages . . . . .	6
4.2	Data Pre-processing . . . . .	6
4.3	Initialization . . . . .	7
4.4	Activation funcion . . . . .	8
4.5	Forward Propagation . . . . .	9
4.6	Cost Function . . . . .	11
4.7	Backward Propagation . . . . .	11
4.8	Update Parameters . . . . .	13
4.9	Merge all together . . . . .	14
<b>5</b>	<b>Train and test the model</b>	<b>15</b>
<b>6</b>	<b>Source code and reference</b>	<b>15</b>

# 1 Problem description

## 1.1 Overview of gender classification

A **gender classification system** uses face of a person from a given image to tell the gender (male/female) of the given person. A successful gender classification approach can boost the performance of many other applications including face recognition and smart human-computer interface.

Usually facial images are used to extract features and then a classifier is applied to the extracted features to learn a gender recognizer. It is an active research topic in Computer Vision and Biometrics fields. The gender classification result is often a binary value, e.g., 1 or 0, representing either male or female. Gender recognition is essentially a two-class classification problem. Although other biometric traits could also be used for gender classification, such as gait, face-based approaches are still the most popular for gender discrimination.

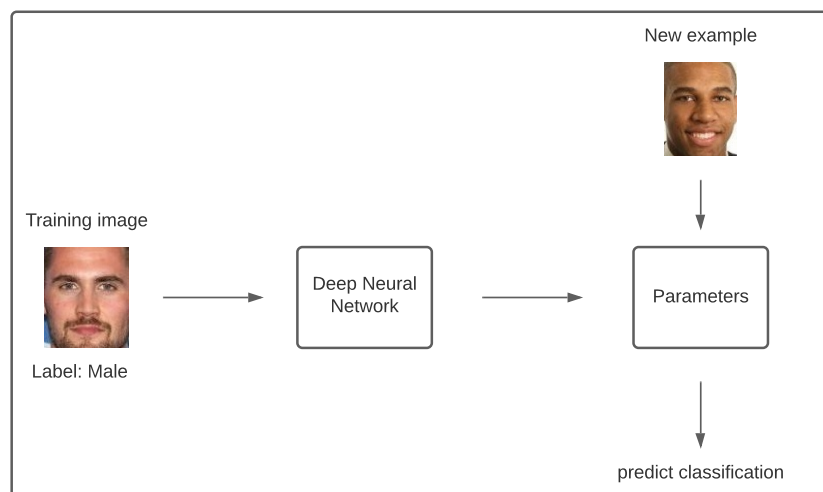


Figure 1: Overview of gender classification

## 1.2 Detailed description

The problem is describe as:

- Get an input image.
- Output the gender prediction: male or female.
- The problem become a binary classification.

## 2 Deep Neural Networks

### 2.1 Deep L-Layer Neural Network

In this assignment, I will use **Deep Neural Network** with 5 layer. The number of units in each layer is  $n^{[1]} = 30, n^{[2]} = 20, n^{[3]} = 7, n^{[4]} = 5, n^{[5]} = 1$  respectively.




Figure 2: Deep neural network

- The input is a (64,64,3) image which is flattened to a vector of size (12288,1).
- The corresponding vector:  $[x_0, x_1, \dots, x_{12287}]^T$  is then multiplied by the weight matrix  $W^{[1]}$  and then added the intercept  $b^{[1]}$ . The result is called the linear unit.
- Next, take the relu of the linear unit. This process could be repeated several times for each  $(W^{[l]}, b^{[l]})$  depending on the model architecture.
- Finally, take the sigmoid of the final linear unit. If it is greater than 0.5, classify it as male, otherwise classify it as female.

### 2.2 Forward Propagation

- Input:  $a^{[l-1]}$
- Output:  $a^{[l]}, cache(z^{[l]})$

The linear forward propagation computed by the following equations:

$$\begin{aligned}z^{[l]} &= w^{[l]}a^{[l-1]} + b[l] \\ a^{[l]} &= g^{[l]}(z^{[l]})\end{aligned}$$

With  $m$  examples, we have:

$$Z^{[l]} = W^{[l]}A^{[l-1]} + b[l]$$

with  $A^{[0]} = X$

$$A^{[l]} = g^{[l]}(Z^{[l]})$$

## 2.3 Backward Propagation

- Input:  $da^{[l]}$
- Output:  $da^{[l-1]}, dW^{[l]}, db^{[l]}$

With 1 training sample, we compute as:

$$dZ^{[l]} = \frac{\partial \mathcal{L}}{\partial Z^{[l]}} = da^{[l]} \times g'^{[l]}(Z^{[l]})$$

$$dW^{[l]} = \frac{\partial \mathcal{J}}{\partial W^{[l]}} = dZ^{[l]}a^{[l-1]T}$$

$$db^{[l]} = \frac{\partial \mathcal{J}}{\partial b^{[l]}} = dZ^{[l]}$$

$$da^{[l-1]} = \frac{\partial \mathcal{J}}{\partial a^{[l-1]}} = W^{[l]T}dZ^{[l]}$$

With  $m$  samples, we compute as:

$$dZ^{[l]} = \frac{\partial \mathcal{L}}{\partial Z^{[l]}} = da^{[l]} \times g'^{[l]}(Z^{[l]})$$

$$dW^{[l]} = \frac{\partial \mathcal{J}}{\partial W^{[l]}} = \frac{1}{m}dZ^{[l]}A^{[l-1]T}$$

$$db^{[l]} = \frac{\partial \mathcal{J}}{\partial b^{[l]}} = \frac{1}{m} \sum_{n=1}^m dZ^{[l]}$$

$$dA^{[l-1]} = \frac{\partial \mathcal{J}}{\partial A^{[l-1]}} = W^{[l]T}dZ^{[l]}$$




Figure 3: Backward propagation calculation

## 2.4 Update Paramters

After backward propagation, I will update parameters as the following equations ( $\alpha$  is the learning rate):

$$W^{[l]} = W^{[l]} - \alpha dW^{[l]}$$

$$b^{[l]} = b^{[l]} - \alpha db^{[l]}$$

## 2.5 Cost Function

Compute the cross-entropy cost  $\mathcal{L}$  using the following formula:

$$\mathcal{L} = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{[L](i)}) + (1 - y^{(i)}) \log(1 - a^{[L](i)}))$$

## 3 Dataset

The dataset I use in this assignment is found on Kaggle.

<https://www.kaggle.com/cashutosh/gender-classification-dataset>

This dataset contain about 28500 images of each class (male and female). I just use 149 first images of each class for training and 60 last images for testing.

## 4 Implementation

### 4.1 Packages

First, import all the packages I will need during this assignment.

- **numpy** is the main package for scientific computing with Python.
- **matplotlib** is a library to plot graphs in Python.
- **cv2** is used for resizing the image to 64x64

```
1 import time
2 import scipy
3 from scipy import ndimage
4 import pathlib
5 import numpy as np
6 from numpy import asarray
7 import matplotlib.pyplot as plt
8 from PIL import Image
9 import random
10 import cv2
11 import os
```

### 4.2 Data Pre-processing

Because the size of images in this dataset is not consistent, I have to resize all of images into 64x64

```
1 image_path = 'C:/Users/Admin/PycharmProjects/ML_asm/dataset'
2 i = 0
3 for path in os.listdir(image_path):
4     full_path = os.path.join(image_path, path)
5     img = cv2.imread(full_path)
6     res = cv2.resize(img, dsize=(64,64), interpolation=cv2.INTER_CUBIC)
7     cv2.imwrite('Image' + str(i) + '.jpg', res)
8     i+=1
```

Next, I will read all images and save into two lists (training list and testing list)

```
1 train_images = list()
2 test_images = list()
3
4 image_path = 'C:/Users/Admin/PycharmProjects/ML_asm/newMale'
5 for path in os.listdir(image_path):
6     full_path = os.path.join(image_path, path)
7     img = Image.open(full_path)
8     train_images.append((img, 1))
9
10 image_path = 'C:/Users/Admin/PycharmProjects/ML_asm/newFemale'
11 for path in os.listdir(image_path):
12     full_path = os.path.join(image_path, path)
13     img = Image.open(full_path)
14     train_images.append((img, 0))
15
16 image_path = 'C:/Users/Admin/PycharmProjects/ML_asm/newMale_validation'
17 for path in os.listdir(image_path):
18     full_path = os.path.join(image_path, path)
```

```
19     img = Image.open(full_path)
20     test_images.append((img, 1))
21
22 image_path = 'C:/Users/Admin/PycharmProjects/ML_asm/newFemale_validation'
23 for path in os.listdir(image_path):
24     full_path = os.path.join(image_path, path)
25     img = Image.open(full_path)
26     test_images.append((img, 0))
27
28 print(len(train_images))
29 random.shuffle(train_images)
30 random.shuffle(test_images)
```

Finally, convert images into array and modify the shape

```
1 def load_data():
2     train_data = [asarray(pair[0]) for pair in train_images]
3     test_data = [asarray(pair[0]) for pair in test_images]
4     train_label = [pair[1] for pair in train_images]
5     test_label = [pair[1] for pair in test_images]
6
7     print(np.shape(train_data))
8
9     train_set_x_orig = np.array(train_data) # your train set features
10    train_set_y_orig = np.array(train_label) # your train set labels
11
12    print(np.shape(train_set_x_orig))
13
14    test_set_x_orig = np.array(test_data) # your test set features
15    test_set_y_orig = np.array(test_label) # your test set labels
16
17    classes = np.array((b'Female', b'Male')) # the list of classes
18
19    train_set_y_orig = train_set_y_orig.reshape((1, train_set_y_orig.shape[0]))
20    test_set_y_orig = test_set_y_orig.reshape((1, test_set_y_orig.shape[0]))
21
22    return train_set_x_orig, train_set_y_orig, test_set_x_orig, test_set_y_orig,
23           classes
24
25 train_x_orig, train_y, test_x_orig, test_y, classes = load_data()
26 # Reshape the training and test examples
27 train_x_flatten = train_x_orig.reshape(train_x_orig.shape[0], -1).T
28 test_x_flatten = test_x_orig.reshape(test_x_orig.shape[0], -1).T
29
30 # Standardize data to have feature values between 0 and 1.
31 train_x = train_x_flatten/255.
32 test_x = test_x_flatten/255.
33
34 print ("train_x's shape: " + str(train_x.shape))
35 print ("test_x's shape: " + str(test_x.shape))
```

## 4.3 Initialization

- The model's structure is [LINEAR -> RELU]  $\times$  (L-1) -> LINEAR -> SIGMOID. I.e., it has  $L - 1$  layers using a ReLU activation function followed by an output layer with a sigmoid activation function.
- I will store  $n^{[l]}$ , the number of units in different layers, in a variable *layer\_dims*



```
1 def initialize_parameters_deep(layer_dims):
2     """
3     Arguments:
4     layer_dims -- python array (list) containing the dimensions of each layer in
                    our network
5
6     Returns:
7     parameters -- python dictionary containing your parameters "W1", "b1", ..., "
                    WL", "bL":
8
9                     W1 -- weight matrix of shape (layer_dims[1], layer_dims[1-1])
10                    b1 -- bias vector of shape (layer_dims[1], 1)
11
12     """
13     np.random.seed(3)
14     parameters = {}
15     L = len(layer_dims) # number of layers in the network
16
17     for l in range(1, L):
18         parameters['W' + str(l)] = np.random.randn(layer_dims[l], layer_dims[l-1])
19         * 0.01
20         parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))
21
22         assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l -
23         1]))
24         assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))
25
26     return parameters
```

## 4.4 Activation funcion

In this assignment, I will use two activation functions: **sigmoid** and **relu**.

```
1 def sigmoid(Z):
2     """
3     Implements the sigmoid activation in numpy
4
5     Arguments:
6     Z -- numpy array of any shape
7
8     Returns:
9     A -- output of sigmoid(z), same shape as Z
10    cache -- returns Z as well, useful during backpropagation
11    """
12
13    A = 1/(1+np.exp(-Z))
14    cache = Z
15
16    return A, cache
17
18 def relu(Z):
19     """
20     Implement the RELU function.
21
22     Arguments:
23     Z -- Output of the linear layer, of any shape
24
25     Returns:
26     A -- Post-activation parameter, of the same shape as Z
27     cache -- a python dictionary containing "A" ; stored for computing the
                backward pass efficiently
```

```
28     """
29
30     A = np.maximum(0,Z)
31
32     assert(A.shape == Z.shape)
33
34     cache = Z
35     return A, cache
36
37
38 def relu_backward(dA, cache):
39     """
40     Implement the backward propagation for a single RELU unit.
41
42     Arguments:
43     dA -- post-activation gradient, of any shape
44     cache -- 'Z' where we store for computing backward propagation efficiently
45
46     Returns:
47     dZ -- Gradient of the cost with respect to Z
48     """
49
50     Z = cache
51     dZ = np.array(dA, copy=True) # just converting dz to a correct object.
52
53     # When z <= 0, you should set dz to 0 as well.
54     dZ[Z <= 0] = 0
55
56     assert (dZ.shape == Z.shape)
57
58     return dZ
59
60 def sigmoid_backward(dA, cache):
61     """
62     Implement the backward propagation for a single SIGMOID unit.
63
64     Arguments:
65     dA -- post-activation gradient, of any shape
66     cache -- 'Z' where we store for computing backward propagation efficiently
67
68     Returns:
69     dZ -- Gradient of the cost with respect to Z
70     """
71
72     Z = cache
73
74     s = 1/(1+np.exp(-Z))
75     dZ = dA * s * (1-s)
76
77     assert (dZ.shape == Z.shape)
78
79     return dZ
```

## 4.5 Forward Propagation

First, I will build the linear part of forward propagation.

```
1 def linear_forward(A, W, b):
2     """
3     Implement the linear part of a layer's forward propagation.
```

```
4
5     Arguments:
6     A -- activations from previous layer (or input data): (size of previous layer,
7         number of examples)
8     W -- weights matrix: numpy array of shape (size of current layer, size of
9         previous layer)
10    b -- bias vector, numpy array of shape (size of the current layer, 1)
11
12    Returns:
13    Z -- the input of the activation function, also called pre-activation
14        parameter
15    cache -- a python tuple containing "A", "W" and "b" ; stored for computing the
16        backward pass efficiently
17    """
18
19    Z = np.dot(W, A) + b
20    cache = (A, W, b)
21
22    return Z, cache
```

Next, I will implement forward propagation with activation function.

```
1 def linear_activation_forward(A_prev, W, b, activation):
2     """
3     Implement the forward propagation for the LINEAR->ACTIVATION layer
4
5     Arguments:
6     A_prev -- activations from previous layer (or input data): (size of previous
7         layer, number of examples)
8     W -- weights matrix: numpy array of shape (size of current layer, size of
9         previous layer)
10    b -- bias vector, numpy array of shape (size of the current layer, 1)
11    activation -- the activation to be used in this layer, stored as a text string
12        : "sigmoid" or "relu"
13
14    Returns:
15    A -- the output of the activation function, also called the post-activation
16        value
17    cache -- a python tuple containing "linear_cache" and "activation_cache";
18        stored for computing the backward pass efficiently
19    """
20
21    if activation == "sigmoid":
22        Z, linear_cache = linear_forward(A_prev, W, b)
23        A, activation_cache = sigmoid(Z)
24
25    elif activation == "relu":
26        Z, linear_cache = linear_forward(A_prev, W, b)
27        A, activation_cache = relu(Z)
28
29    cache = (linear_cache, activation_cache)
30
31    return A, cache
```

Finally, the full version of forward propagation

```
1 def L_model_forward(X, parameters):
2     """
3     Implement forward propagation for the [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID
4         computation
5
6     Arguments:
7     X -- data, numpy array of shape (input size, number of examples)
```

```
7     parameters -- output of initialize_parameters_deep()
8
9     Returns:
10    AL -- activation value from the output (last) layer
11    caches -- list of caches containing:
12              every cache of linear_activation_forward() (there are L of them,
13              indexed from 0 to L-1)
14    """
15
16    caches = []
17    A = X
18    L = len(parameters) // 2    # number of layers in the neural network
19
20    for l in range(1, L):
21        A, cache = linear_activation_forward(A_prev, parameters["W" + str(l)],
22        parameters["b" + str(l)], activation = "relu")
23
24        caches.append(cache)
25
26    AL, cache = linear_activation_forward(A, parameters["W"+ str(L)], parameters["
27    b" + str(L)], activation = "sigmoid")
28    caches.append(cache)
29
30    return AL, caches
```

## 4.6 Cost Function

```
1 def compute_cost(AL, Y):
2     """
3     Implement the cost function defined by equation (7).
4
5     Arguments:
6     AL -- probability vector corresponding to your label predictions, shape (1,
7     number of examples)
8     Y -- true "label" vector (for example: containing 0 if non-cat, 1 if cat),
9     shape (1, number of examples)
10
11     Returns:
12     cost -- cross-entropy cost
13     """
14
15     m = Y.shape[1]
16
17     cost = -np.sum(np.dot(Y, np.log(AL).T) + np.dot(1 - Y, np.log(1 - AL).T)) / m
18
19     cost = np.squeeze(cost)
20
21     return cost
```

## 4.7 Backward Propagation

Backpropagation is used to calculate the gradient of the loss function with respect to the parameters.

```
1 def linear_backward(dZ, cache):
```

```
2 """
3 Implement the linear portion of backward propagation for a single layer (layer
4 1)
5
6 Arguments:
7 dZ -- Gradient of the cost with respect to the linear output (of current layer
8 1)
9 cache -- tuple of values (A_prev, W, b) coming from the forward propagation in
10 the current layer
11
12 Returns:
13 dA_prev -- Gradient of the cost with respect to the activation (of the
14 previous layer l-1), same shape as A_prev
15 dW -- Gradient of the cost with respect to W (current layer 1), same shape as
16 W
17 db -- Gradient of the cost with respect to b (current layer 1), same shape as
18 b
19 """
20 A_prev, W, b = cache
21 m = A_prev.shape[1]
22
23 dW = np.dot(dZ, A_prev.T) / m
24 db = np.sum(dZ, axis = 1, keepdims = True) / m
25 dA_prev = np.dot(W.T, dZ)
26
27 return dA_prev, dW, db
```

Next, I will implement backward propagation with activation function.

```
1 def linear_activation_backward(dA, cache, activation):
2     """
3     Implement the backward propagation for the LINEAR->ACTIVATION layer.
4
5     Arguments:
6     dA -- post-activation gradient for current layer l
7     cache -- tuple of values (linear_cache, activation_cache) we store for
8     computing backward propagation efficiently
9     activation -- the activation to be used in this layer, stored as a text string
10     : "sigmoid" or "relu"
11
12     Returns:
13     dA_prev -- Gradient of the cost with respect to the activation (of the
14     previous layer l-1), same shape as A_prev
15     dW -- Gradient of the cost with respect to W (current layer l), same shape as
16     W
17     db -- Gradient of the cost with respect to b (current layer l), same shape as
18     b
19     """
20     linear_cache, activation_cache = cache
21
22     if activation == "relu":
23         dZ = relu_backward(dA, activation_cache)
24         dA_prev, dW, db = linear_backward(dZ, linear_cache)
25
26     elif activation == "sigmoid":
27         dZ = sigmoid_backward(dA, activation_cache)
28         dA_prev, dW, db = linear_backward(dZ, linear_cache)
29
30     return dA_prev, dW, db
```

Finally, the full version of backward propagation

```
1 def L_model_backward(AL, Y, caches):
```

```
2 """
3 Implement the backward propagation for the [LINEAR->RELU] * (L-1) -> LINEAR ->
  SIGMOID group
4
5 Arguments:
6 AL -- probability vector, output of the forward propagation (L_model_forward())
7 Y -- true "label" vector (containing 0 if non-cat, 1 if cat)
8 caches -- list of caches containing:
9           every cache of linear_activation_forward() with "relu" (it's
10 caches[l], for l in range(L-1) i.e l = 0...L-2)
11           the cache of linear_activation_forward() with "sigmoid" (it's
12 caches[L-1])
13
14 Returns:
15 grads -- A dictionary with the gradients
16           grads["dA" + str(l)] = ...
17           grads["dW" + str(l)] = ...
18           grads["db" + str(l)] = ...
19
20 """
21 grads = {}
22 L = len(caches) # the number of layers
23 m = AL.shape[1]
24 Y = Y.reshape(AL.shape) # after this line, Y is the same shape as AL
25
26 # Initializing the backpropagation
27 dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))
28
29 current_cache = caches[L-1]
30 dA_prev_temp, dW_temp, db_temp = linear_activation_backward(dAL, current_cache
31 , activation = "sigmoid")
32 grads["dA" + str(L-1)] = dA_prev_temp
33 grads["dW" + str(L)] = dW_temp
34 grads["db" + str(L)] = db_temp
35
36 # Loop from l=L-2 to l=0
37 for l in reversed(range(L-1)):
38     current_cache = caches[l]
39     dA_prev_temp, dW_temp, db_temp = linear_activation_backward(dA_prev_temp,
40 current_cache, activation = "relu")
41     grads["dA" + str(l)] = dA_prev_temp
42     grads["dW" + str(l + 1)] = dW_temp
43     grads["db" + str(l + 1)] = db_temp
44
45 return grads
```

## 4.8 Update Parameters

```
1 def update_parameters(params, grads, learning_rate):
2     """
3     Update parameters using gradient descent
4
5     Arguments:
6     params -- python dictionary containing your parameters
7     grads -- python dictionary containing your gradients, output of
8     L_model_backward
```

```
9     Returns:
10     parameters -- python dictionary containing your updated parameters
11                 parameters["W" + str(l)] = ...
12                 parameters["b" + str(l)] = ...
13
14     parameters = params.copy()
15     L = len(parameters) // 2 # number of layers in the neural network
16     print(grads)
17
18     for l in range(L):
19         parameters['W' + str(l+1)] = parameters['W' + str(l+1)] - learning_rate*
grads["dW" + str(l+1)]
20         parameters["b" + str(l+1)] = parameters["b" + str(l+1)] - learning_rate*
grads["db" + str(l+1)]
21
22     return parameters
```

## 4.9 Merge all together

```
1 def L_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations =
2     3000, print_cost=False):
3     """
4     Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID.
5
6     Arguments:
7     X -- data, numpy array of shape (num_px * num_px * 3, number of examples)
8     Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of shape (1,
9     number of examples)
10    layers_dims -- list containing the input size and each layer size, of length (
11    number of layers + 1).
12    learning_rate -- learning rate of the gradient descent update rule
13    num_iterations -- number of iterations of the optimization loop
14    print_cost -- if True, it prints the cost every 100 steps
15
16    Returns:
17    parameters -- parameters learnt by the model. They can then be used to predict
18    .
19    """
20
21    np.random.seed(1)
22    costs = [] # keep track of cost
23
24    parameters = initialize_parameters_deep(layers_dims)
25
26    # Loop (gradient descent)
27    for i in range(0, num_iterations):
28        AL, caches = L_model_forward(X, parameters)
29        cost = compute_cost(AL, Y)
30        grads = L_model_backward(AL, Y, caches)
31        parameters = update_parameters(parameters, grads, learning_rate)
32
33        # Print the cost every 100 iterations
34        if print_cost and i % 100 == 0 or i == num_iterations - 1:
35            print("Cost after iteration {}: {}".format(i, np.squeeze(cost)))
36        if i % 100 == 0 or i == num_iterations:
37            costs.append(cost)
38
39    return parameters, costs
```

## 5 Train and test the model

I will train the model with 2500 epochs

```
1 layers_dims = [12288, 30, 20, 7, 5, 1] # 5-layer model
2 parameters, costs = L_layer_model(train_x, train_y, layers_dims, num_iterations =
    2500, print_cost = True)
```

The result displayed on screen:

```
Cost after iteration 0: 0.6954910520868507
Cost after iteration 100: 0.6363603599908855
Cost after iteration 200: 0.5951576846762471
Cost after iteration 300: 0.5356822577317358
Cost after iteration 400: 0.6437753935550212
Cost after iteration 500: 0.6005912827800174
Cost after iteration 600: 0.41064940228175373
Cost after iteration 700: 0.4713705382972464
Cost after iteration 800: 0.45382283383380806
Cost after iteration 900: 0.5756933350193408
Cost after iteration 1000: 0.43581980393225767
Cost after iteration 1100: 0.25193893442340576
Cost after iteration 1200: 0.6394347638282675
Cost after iteration 1300: 0.8704262146564183
Cost after iteration 1400: 0.6009575015151132
Cost after iteration 1500: 0.3736721004281588
Cost after iteration 1600: 0.21695661007327502
Cost after iteration 1700: 0.6208097949566367
Cost after iteration 1800: 0.08488470674478711
Cost after iteration 1900: 0.172775495132184
Cost after iteration 2000: 1.329117758604708
Cost after iteration 2100: 0.27716715465070896
Cost after iteration 2200: 0.056080448570820104
Cost after iteration 2300: 0.29480557061918305
Cost after iteration 2400: 0.014825415876161705
Cost after iteration 2499: 0.007212752044204772
```

```
1 pred_train = predict(train_x, train_y, parameters)
```

Accuracy: 0.9999999999999998

```
1 pred_test = predict(test_x, test_y, parameters)
```

Accuracy: 0.9333333333333333

## 6 Source code and reference

Source code I upload all the source code and dataset here:

<https://github.com/dinhhoanganh2001/Gender-Classification>

Notebook version:

<https://colab.research.google.com/drive/1KoVMBaBXttQ1Qf2H0XYuwtExxD840Gen?usp=sharing>





## References

- [1] Vũ Hữu Tiếp (2018), Machine Learning cơ bản.
- [2] Neural Networks and Deep Learning (Course on Coursera)-Andrew Ng.
- [3] B.A.Golomb, D.T.Lawrence, T.J.Sejnowski (1991), SEXNET: A Neural Network Identifies Sex from Human Faces.